

Lab Exercise 10- Implementing Resource Quota in Kubernetes

Objective:

In Kubernetes, Resource Quotas are used to control the resource consumption of namespaces. They help in managing and enforcing limits on the usage of resources like CPU, memory, and the number of objects (e.g., Pods, Services) within a namespace. This exercise will guide you through creating and managing Resource Quotas to limit the resources used by applications in a specific namespace.

Step 1: Understand Resource Quotas

Resource Quotas allow you to:

- Limit the amount of CPU and memory a namespace can use.
- Control the number of certain types of resources (e.g., Pods, Services, PersistentVolumeClaims) in a namespace.
- Prevent a namespace from consuming more resources than allocated, ensuring fair usage across multiple teams or applications.

Step 2: Create a Namespace

First, create a namespace where you will apply the Resource Quota. This helps in isolating and controlling resource usage within that specific namespace.

Create a YAML file named ***quota-namespace.yaml*** with the following content:

```
apiVersion: v1
kind: Namespace
metadata:
  name: quota-example # The name of the namespace.
```

```
← → K8S
quota-namespace.yaml ×
quota-namespace.yaml
1  apiVersion: v1
2  kind: Namespace
3  metadata:
4  |   name: quota-example    # The name of the namespace.
```

Apply the YAML to create the namespace:

```
kubectl apply -f quota-namespace.yaml
```

```
[sai@Sais-Mac K8S % kubectl apply -f quota-namespace.yaml
namespace/quota-example created
```

Verify that the namespace is created:

```
kubectl get namespaces
```

```
[sai@Sais-Mac K8S % kubectl get ns
NAME                STATUS    AGE
default             Active    14h
kube-node-lease     Active    14h
kube-public         Active    14h
kube-system         Active    14h
quota-example       Active    31s
```

You should see quota-example listed in the output.

Step 3: Define a Resource Quota

Next, create a Resource Quota YAML file named **resource-quota.yaml** with the following content:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: example-quota  # The name of the Resource Quota.
```

```

namespace: quota-example # The namespace to which the Resource Quota will apply.
spec:
  hard:
    # The hard limits imposed by this Resource Quota.
    requests.cpu: "2" # The total CPU resource requests allowed in the namespace (2 cores).
    requests.memory: "4Gi" # The total memory resource requests allowed in the namespace (4 GiB).
    limits.cpu: "4" # The total CPU resource limits allowed in the namespace (4 cores).
    limits.memory: "8Gi" # The total memory resource limits allowed in the namespace (8 GiB).
    pods: "10" # The total number of Pods allowed in the namespace.
    persistentvolumeclaims: "5" # The total number of PersistentVolumeClaims allowed in the namespace.
    configmaps: "10" # The total number of ConfigMaps allowed in the namespace.
    services: "5" # The total number of Services allowed in the namespace.

```

```

1  apiVersion: v1
2  kind: ResourceQuota
3  metadata:
4    name: example-quota # The name of the Resource Quota.
5    namespace: quota-example # The namespace to which the Resource Quota will apply.
6  spec:
7    hard:
8      # The hard limits imposed by this Resource Quota.
9      requests.cpu: "2" # The total CPU resource requests allowed in the namespace (2 cores).
10     requests.memory: "4Gi" # The total memory resource requests allowed in the namespace (4 GiB).
11     limits.cpu: "4" # The total CPU resource limits allowed in the namespace (4 cores).
12     limits.memory: "8Gi" # The total memory resource limits allowed in the namespace (8 GiB).
13     pods: "10" # The total number of Pods allowed in the namespace.
14     persistentvolumeclaims: "5" # The total number of PersistentVolumeClaims allowed in the namespace.
15     configmaps: "10" # The total number of ConfigMaps allowed in the namespace.
16     services: "5" # The total number of Services allowed in the namespace.

```

Step 4: Apply the Resource Quota

Apply the Resource Quota YAML to the namespace:

```
kubectl apply -f resource-quota.yaml
```

```
[sai@Sais-Mac K8S % kubectl apply -f resource-quota.yaml
resourcequota/example-quota created
```

Verify that the Resource Quota is applied:

```
kubectl get resourcequota -n quota-example
```

```

sai@Sais-Mac K8S % kubectl get resourcequota -n quota-example
NAME          AGE   REQUEST
example-quota 29s   configmaps: 1/10, persistentvolumeclaims: 0/5, pods: 0/10, requests.cpu: 0/2, requests.memory: 0/4Gi, services: 0/5
sai@Sais-Mac K8S %

```

To see the details of the applied Resource Quota:

```
kubectl describe resourcequota example-quota -n quota-example
```

```
sai@Sais-Mac K8S % kubectl describe resourcequota example-quota -n quota-example
Name:          example-quota
Namespace:     quota-example
Resource       Used  Hard
-----
configmaps    1    10
limits.cpu    0     4
limits.memory 0    8Gi
persistentvolumeclaims 0     5
pods          0    10
requests.cpu   0     2
requests.memory 0    4Gi
services      0     5
```

Step 5: Test the Resource Quota

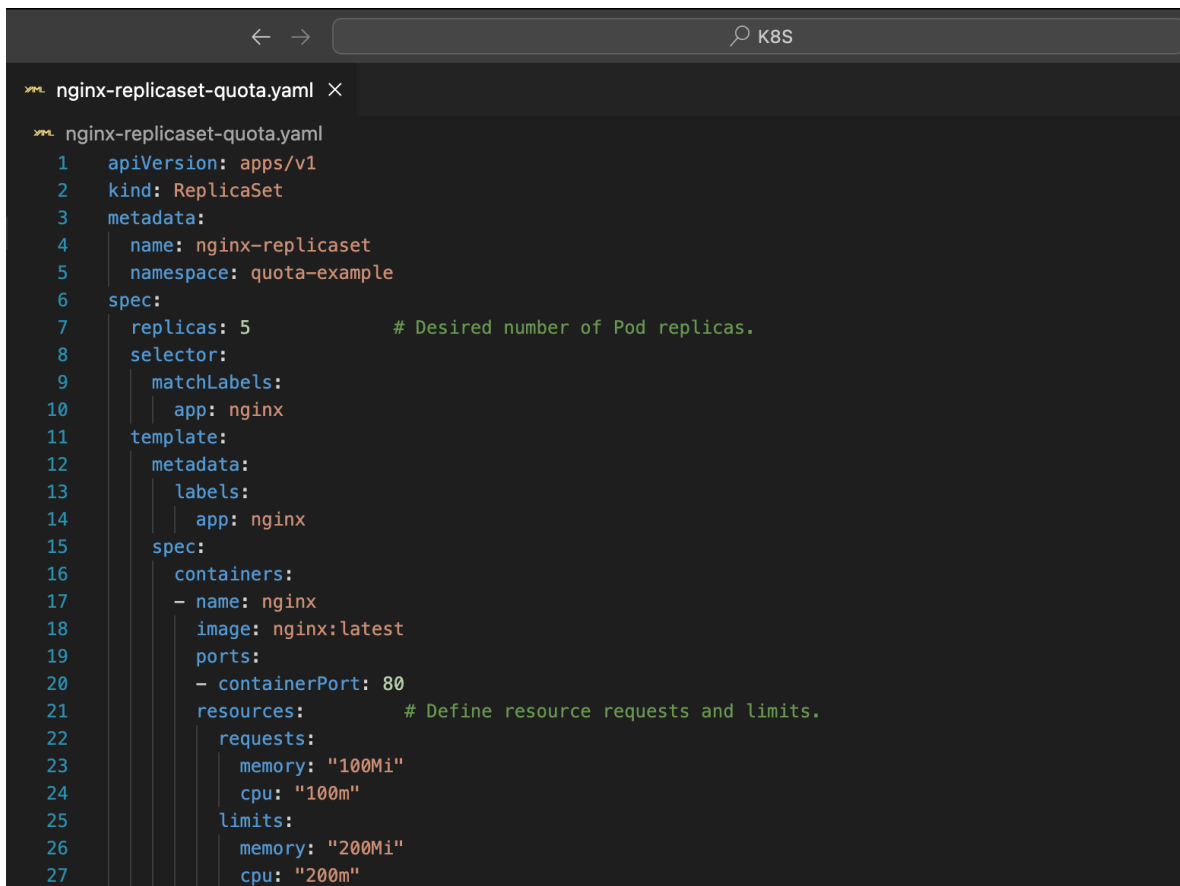
Let's create some resources in the quota-example namespace to see how the Resource Quota affects them.

Deploy a ReplicaSet with Resource Requests and Limits

Create a YAML file named ***nginx-replicaset-quota.yaml*** with the following content:

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  namespace: quota-example
spec:
  replicas: 5      # Desired number of Pod replicas.
  selector:
    matchLabels:
      app: nginx
```

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
    - name: nginx
      image: nginx:latest
      ports:
      - containerPort: 80
    resources:      # Define resource requests and limits.
      requests:
        memory: "100Mi"
        cpu: "100m"
      limits:
        memory: "200Mi"
        cpu: "200m"
```



The screenshot shows a terminal window with a dark background and light-colored text. At the top, there is a navigation bar with a back arrow, a forward arrow, and a search bar containing the text "K8S". Below the navigation bar, the terminal displays the contents of a file named "nginx-replicaset-quota.yaml". The file content is as follows:

```
1  apiVersion: apps/v1
2  kind: ReplicaSet
3  metadata:
4    name: nginx-replicaset
5    namespace: quota-example
6  spec:
7    replicas: 5          # Desired number of Pod replicas.
8    selector:
9      matchLabels:
10       app: nginx
11  template:
12    metadata:
13      labels:
14       app: nginx
15    spec:
16      containers:
17      - name: nginx
18        image: nginx:latest
19        ports:
20        - containerPort: 80
21      resources:      # Define resource requests and limits.
22        requests:
23          memory: "100Mi"
24          cpu: "100m"
25        limits:
26          memory: "200Mi"
27          cpu: "200m"
```

Explanation:

This ReplicaSet requests a total of 500m CPU and 500Mi memory across 5 replicas. It also limits each replica to use a maximum of 200m CPU and 200Mi memory.

Apply this YAML to create the ReplicaSet:

```
kubectl apply -f nginx-replicaset-quota.yaml
```

```
sai@Sais-Mac K8S % kubectl apply -f nginx-replicaset-quota.yaml
replicaset.apps/nginx-replicaset created
```

Check the status of the Pods and ensure they are created within the constraints of the Resource Quota:

```
kubectl get pods -n quota-example
```

```
sai@Sais-Mac K8S % kubectl get pods -n quota-example
NAME                                READY   STATUS    RESTARTS   AGE
nginx-replicaset-6k4mb             1/1     Running   0           37s
nginx-replicaset-cp6sf             1/1     Running   0           37s
nginx-replicaset-fmx8c             1/1     Running   0           37s
nginx-replicaset-lg6wk             1/1     Running   0           37s
nginx-replicaset-tbnkf             1/1     Running   0           37s
sai@Sais-Mac K8S %
```

To describe the Pods and see their resource allocations:

```
kubectl describe pods -l app=nginx -n quota-example
```

Attempt to Exceed the Resource Quota

```

[sai@Sais-Mac K8S % kubectl describe pods -l app=nginx -n quota-example
Name:          nginx-replicaset-6k4mb
Namespace:     quota-example
Priority:       0
Service Account: default
Node:          docker-desktop/192.168.65.3
Start Time:    Tue, 12 Nov 2024 01:34:54 +0530
Labels:        app=nginx
Annotations:    <none>
Status:        Running
IP:            10.1.0.19
IPs:
  IP:          10.1.0.19
Controlled By: ReplicaSet/nginx-replicaset
Containers:
  nginx:
    Container ID:  docker://93a6151480de03309483fb4047e80cd667034800294e555f86202effbeda1337
    Image:         nginx:latest
    Image ID:      docker-pullable://nginx@sha256:28402db69fec7c17e179ea87882667f1e054391138f77ffaf0c3eb388efc3fffb
    Port:          80/TCP
    Host Port:     0/TCP
    State:         Running
      Started:     Tue, 12 Nov 2024 01:35:06 +0530
    Ready:         True
    Restart Count:  0
    Limits:
      cpu:         200m
      memory:      200Mi
    Requests:
      cpu:         100m
      memory:      100Mi
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-wqdn8 (ro)
Conditions:
  Type                 Status
  PodReadyToStartContainers  True
  Initialized            True
  Ready                  True
  ContainersReady         True
  PodScheduled           True
Volumes:
  kube-api-access-wqdn8:
    Type:              Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:      kube-root-ca.crt
    ConfigMapOptional:  <nil>
    DownwardAPI:        true
  QoS Class:           Burstable
  Node-Selectors:      <none>
  Tolerations:         node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                      node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   106s  default-scheduler  Successfully assigned quota-example/nginx-replicaset-6k4mb to docker-desktop
  Normal   Pulling     105s  kubelet         Pulling image "nginx:latest"
  Normal   Pulled      94s   kubelet         Successfully pulled image "nginx:latest" in 3.241s (11.083s including waiting). Image size: 196880357 bytes.
  Normal   Created     94s   kubelet         Created container nginx
  Normal   Started     94s   kubelet         Started container nginx

Name:          nginx-replicaset-cp6sf
Namespace:     quota-example
Priority:       0

```

Try creating additional resources to see if they are rejected when exceeding the quota. For example, create more Pods or increase the CPU/memory requests to exceed the quota limits.

Create a YAML file named ***nginx-extra-pod.yaml*** with the following content:

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx-extra-pod
```

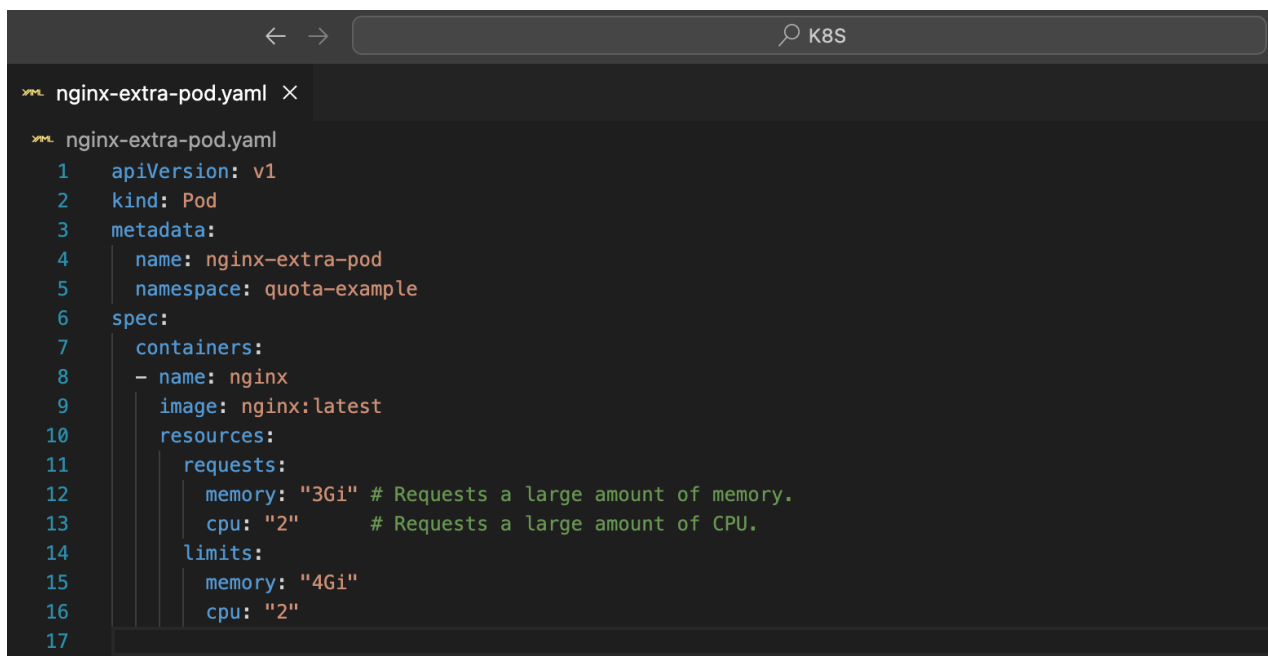
```
  namespace: quota-example
```

```
spec:
```

```
  containers:
```

```
  - name: nginx
```

```
image: nginx:latest
resources:
  requests:
    memory: "3Gi" # Requests a large amount of memory.
    cpu: "2"      # Requests a large amount of CPU.
  limits:
    memory: "4Gi"
    cpu: "2"
```

A screenshot of a code editor window titled 'nginx-extra-pod.yaml'. The editor shows the following YAML content:

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: nginx-extra-pod
5    namespace: quota-example
6  spec:
7    containers:
8    - name: nginx
9      image: nginx:latest
10     resources:
11       requests:
12         memory: "3Gi" # Requests a large amount of memory.
13         cpu: "2"      # Requests a large amount of CPU.
14       limits:
15         memory: "4Gi"
16         cpu: "2"
```

Apply this YAML to create the Pod:

```
kubectl apply -f nginx-extra-pod.yaml
```

```
sai@Sais-Mac K8S % kubectl apply -f nginx-extra-pod.yaml
Error from server (Forbidden): error when creating "nginx-extra-pod.yaml": pods
"nginx-extra-pod" is forbidden: exceeded quota: example-quota, requested: requests.cpu=2, used: requests.cpu=500m, limited: requests.cpu=2
sai@Sais-Mac K8S %
```


This should fail due to exceeding the Resource Quota. Check the events to see the failure reason:

```
kubectl get events -n quota-example
```

```
sai@Sais-Mac K8S % kubectl get events --n quota-example
LAST SEEN   TYPE      REASON      OBJECT                                MESSAGE
5m17s       Normal    Scheduled    pod/nginx-replicaset-6k4mb           Successfully assigned quota-example/nginx-replicaset-6k4mb to docker-desktop
5m16s       Normal    Pulling      pod/nginx-replicaset-6k4mb           Pulling image "nginx:latest"
5m5s        Normal    Pulled       pod/nginx-replicaset-6k4mb           Successfully pulled image "nginx:latest" in 3.241s (11.083s including waiting). Image size: 196880357 bytes.
5m5s        Normal    Created      pod/nginx-replicaset-6k4mb           Created container nginx
5m17s       Normal    Scheduled    pod/nginx-replicaset-cp6sf           Successfully assigned quota-example/nginx-replicaset-cp6sf to docker-desktop
5m16s       Normal    Pulling      pod/nginx-replicaset-cp6sf           Pulling image "nginx:latest"
5m13s       Normal    Pulled       pod/nginx-replicaset-cp6sf           Successfully pulled image "nginx:latest" in 3.267s (3.268s including waiting). Image size: 196880357 bytes.
5m13s       Normal    Created      pod/nginx-replicaset-cp6sf           Created container nginx
5m13s       Normal    Started      pod/nginx-replicaset-cp6sf           Started container nginx
5m17s       Normal    Scheduled    pod/nginx-replicaset-fmx8c           Successfully assigned quota-example/nginx-replicaset-fmx8c to docker-desktop
5m16s       Normal    Pulling      pod/nginx-replicaset-fmx8c           Pulling image "nginx:latest"
5m11s       Normal    Pulled       pod/nginx-replicaset-fmx8c           Successfully pulled image "nginx:latest" in 2.536s (5.804s including waiting). Image size: 196880357 bytes.
5m11s       Normal    Created      pod/nginx-replicaset-fmx8c           Created container nginx
5m11s       Normal    Started      pod/nginx-replicaset-fmx8c           Started container nginx
5m17s       Normal    Scheduled    pod/nginx-replicaset-lg6wk           Successfully assigned quota-example/nginx-replicaset-lg6wk to docker-desktop
5m16s       Normal    Pulling      pod/nginx-replicaset-lg6wk           Pulling image "nginx:latest"
5m2s        Normal    Pulled       pod/nginx-replicaset-lg6wk           Successfully pulled image "nginx:latest" in 2.886s (13.969s including waiting). Image size: 196880357 bytes.
5m2s        Normal    Created      pod/nginx-replicaset-lg6wk           Created container nginx
5m2s        Normal    Started      pod/nginx-replicaset-lg6wk           Started container nginx
5m17s       Normal    Scheduled    pod/nginx-replicaset-tbnkf           Successfully assigned quota-example/nginx-replicaset-tbnkf to docker-desktop
5m16s       Normal    Pulling      pod/nginx-replicaset-tbnkf           Pulling image "nginx:latest"
5m8s        Normal    Pulled       pod/nginx-replicaset-tbnkf           Successfully pulled image "nginx:latest" in 2.299s (7.841s including waiting). Image size: 196880357 bytes.
5m8s        Normal    Created      pod/nginx-replicaset-tbnkf           Created container nginx
5m8s        Normal    Started      pod/nginx-replicaset-tbnkf           Started container nginx
5m17s       Normal    SuccessfulCreate replicaset/nginx-replicaset          Created pod: nginx-replicaset-6k4mb
5m17s       Normal    SuccessfulCreate replicaset/nginx-replicaset          Created pod: nginx-replicaset-lg6wk
5m17s       Normal    SuccessfulCreate replicaset/nginx-replicaset          Created pod: nginx-replicaset-tbnkf
5m17s       Normal    SuccessfulCreate replicaset/nginx-replicaset          Created pod: nginx-replicaset-cp6sf
5m17s       Normal    SuccessfulCreate replicaset/nginx-replicaset          Created pod: nginx-replicaset-fmx8c
sai@Sais-Mac K8S %
```

Look for error messages indicating that the Pod creation was denied due to resource constraints.

Step 6: Clean Up Resources

To delete the resources you created:

```
kubectl delete -f nginx-replicaset-quota.yaml
```

```
kubectl delete -f nginx-extra-pod.yaml
```

```
kubectl delete -f resource-quota.yaml
```

```
kubectl delete namespace quota-example
```

```
sai@Sais-Mac K8S % kubectl delete -f nginx-replicaset-quota.yaml
kubectl delete -f nginx-extra-pod.yaml
kubectl delete -f resource-quota.yaml
[kubectl delete namespace quota-example
replicaset.apps "nginx-replicaset" deleted
Error from server (NotFound): error when deleting "nginx-extra-pod.yaml": pods "
nginx-extra-pod" not found
resourcequota "example-quota" deleted
namespace "quota-example" deleted
sai@Sais-Mac K8S %
```